

# Project4

---

Multiple Column & Join Table

# Before doing project 4..

- We allow you to use C++ features, not only C.
- Whether you use C or C++, you need to modify the name of existing delete API.
  - ~~int delete~~ (int table\_id, int64\_t key);
  - **int remove** (int table\_id, int64\_t key);

# Multiple column

---

- Your database only supports storing a key and a value in a single record.
- Our first goal is to implement a multiple column record.

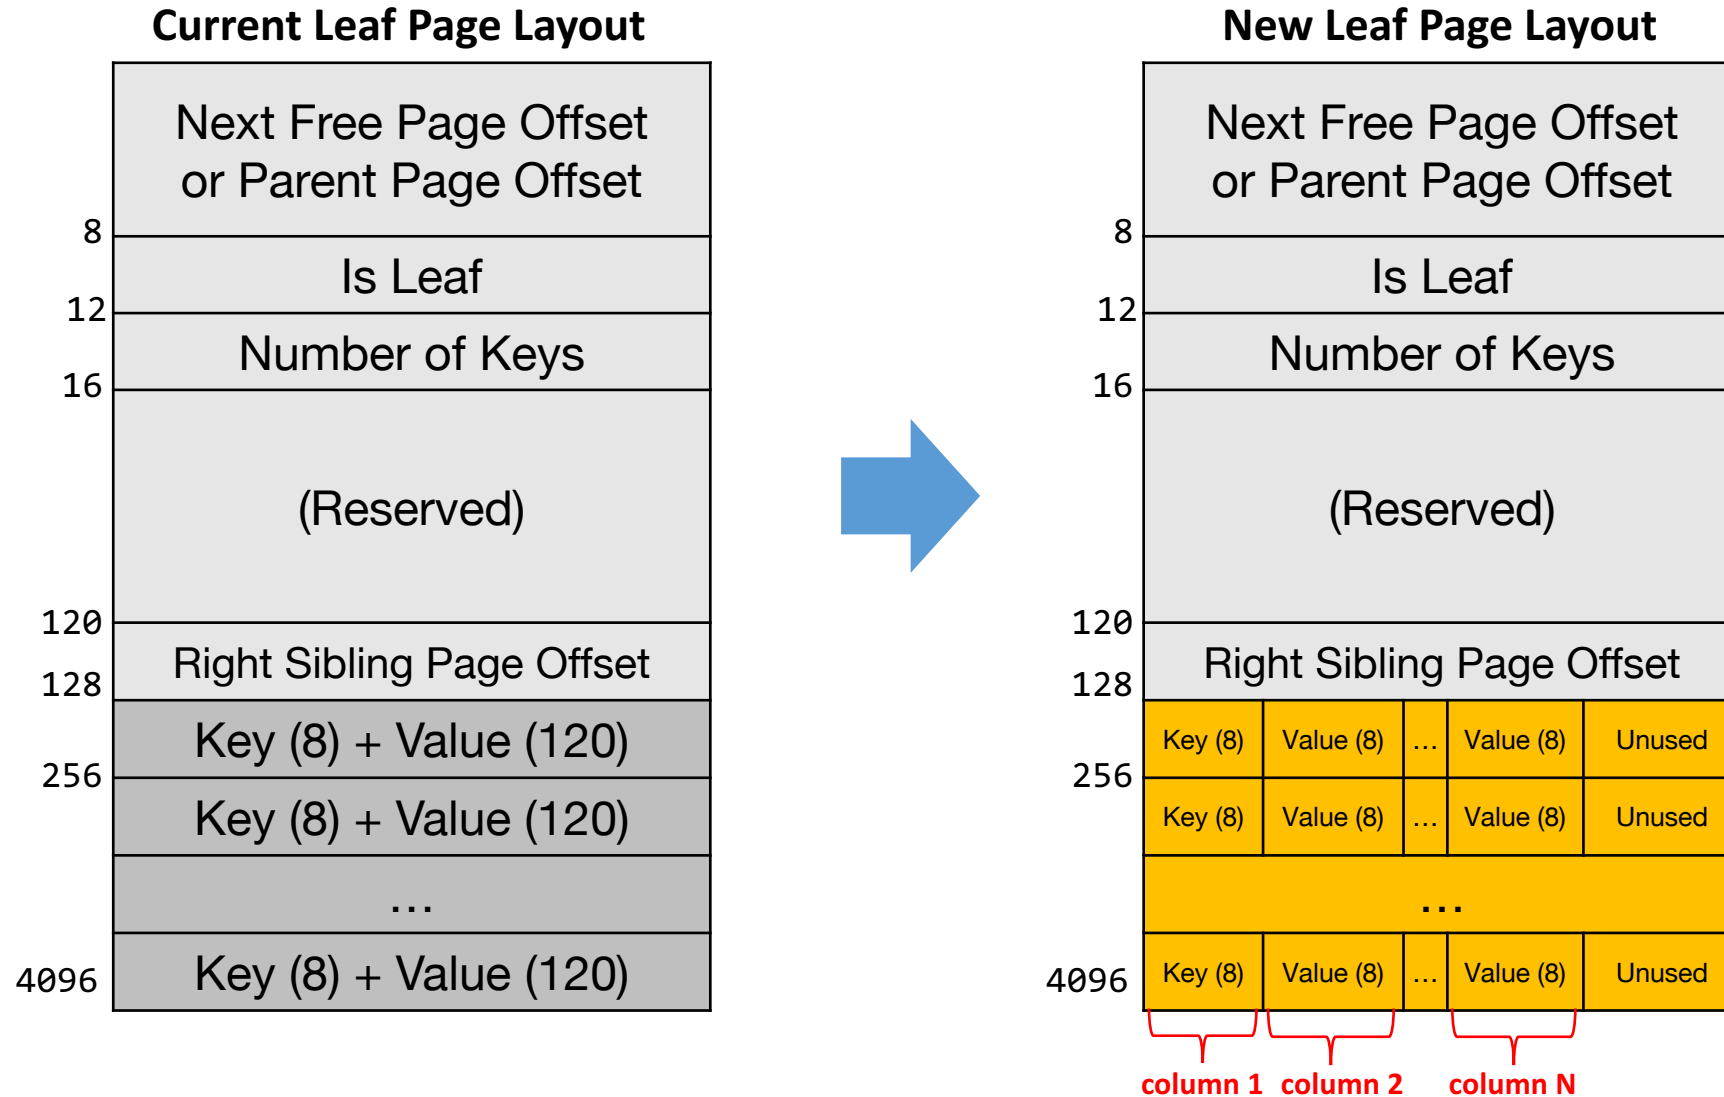
# Need to change APIs

1. `int open_table (char * pathname, int num_column);`
  - If 'pathname' data file exist, open it and ignore the num\_column parameter.
  - Otherwise, create a new one with **num\_column** columns.
  - $2 \leq \text{num\_column} \leq 16$ 
    - (We suppose there must a key and at least one value in a single record)
2. `int insert (int table_id, int64_t key, int64_t* values);`
  - Insert a record with given key and **an array of values**
  - Example

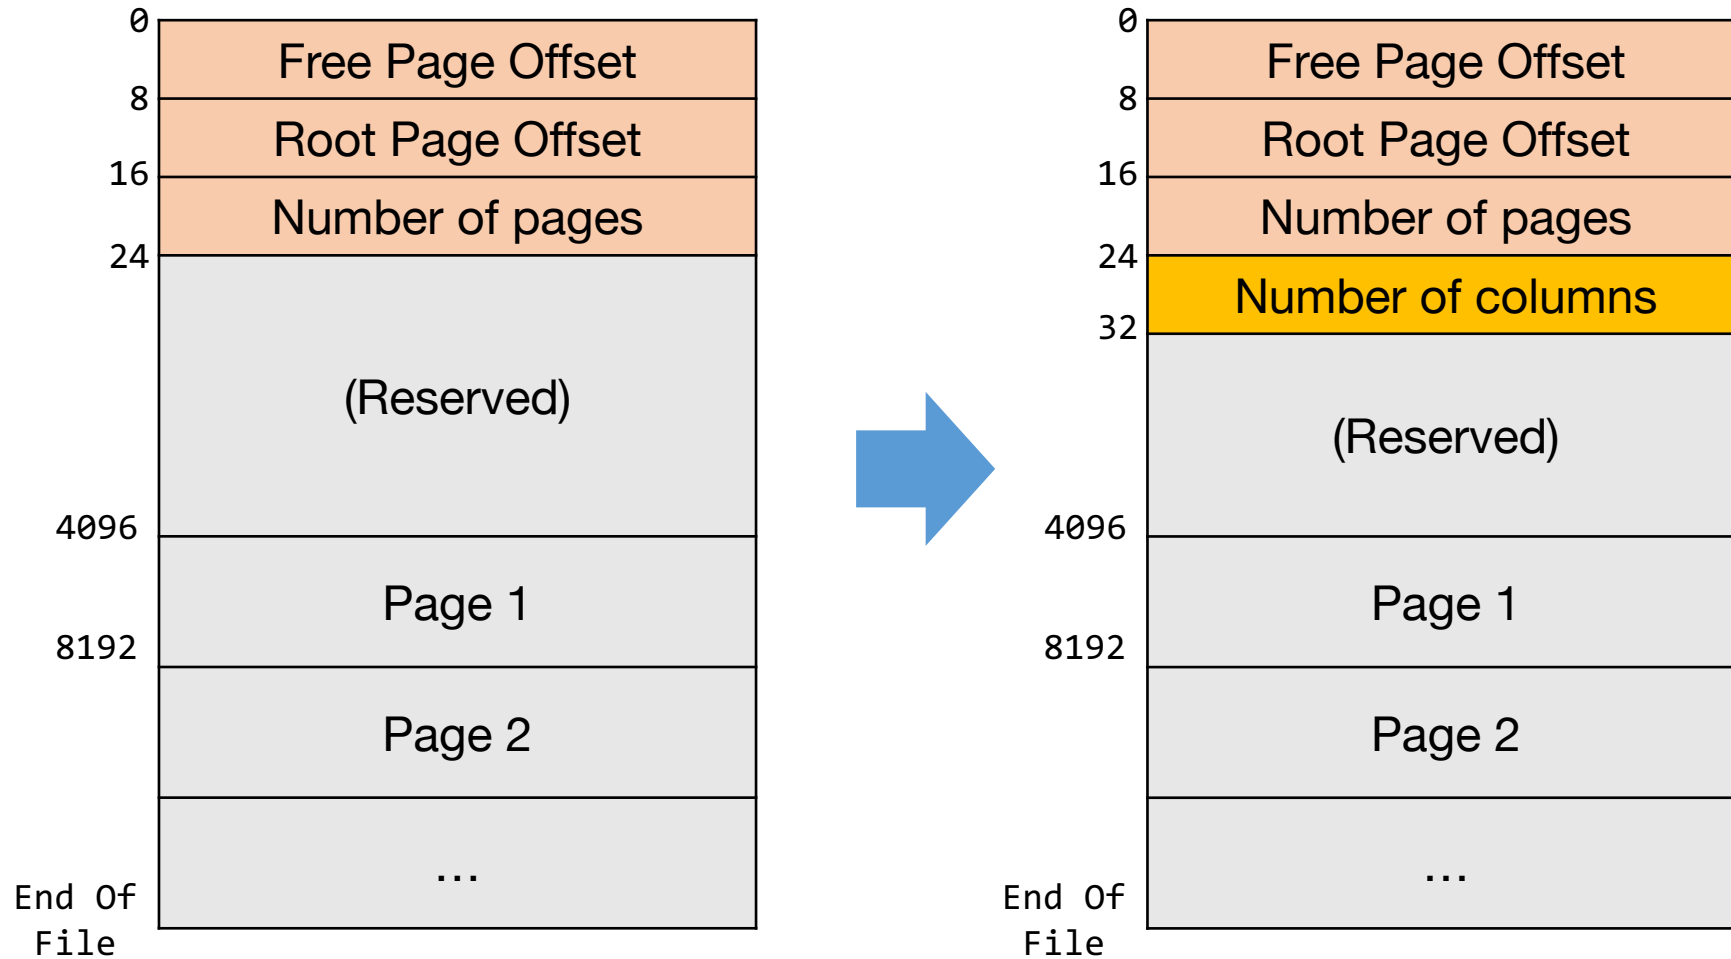
```
int64_t values = (int64_t*)malloc(sizeof(int64_t) * (ncolumn - 1));  
for (int i = 0; i < ncolumn - 1; i++) values[i] = 1234;  
int ret = insert(tid, key, values);
```

3. `int64_t * find (int table_id, int64_t key);`
  - Returns a pointer of an array of values(do not include a key) if the key is found.

# New leaf page format



# New header page format



# Join Tables

- Your database system doesn't consider JOIN table operation yet.
- Our second goal is to implement a **JOIN query** that consist of multiple join operations and an aggregation.

# Join Tables


- You need to support this API
  - **int64\_t join (const char\* query);**
  - *query* format
    - “ $T_1.C_1=T_2.C_2$ ” in case the number of join op. = 1
    - “ $T_1.C_1=T_2.C_2\&T_3.C_3=T_4.C_4$ ” in case the number of join op. = 2
    - “ $T_1.C_1=T_2.C_2\&T_3.C_3=T_4.C_4\&T_5.C_5=T_6.C_6$ ” in case the number of join op. = 3
    - Same pattern for the number of join op. > 3
  - Return a sum of all keys in the query result.
    - If there is no result, return 0
  - $1 \leq \text{number of join op.} \leq 9$



# Join Tables

- Example

- `int64_t ret = join("1.1=2.2&2.3=3.2"); // ret: 29(1+3+9+3+4+9)`




# Join Tables


- Example
  - All join graphs are **connected**. No cross product.
  - All graphs are **acyclic** and there will be **no self join**.
    - [cyclic graph example: "1.1=2.1&2.2=1.2"] [self join example: "1.1=1.2"]
  - You can select a query plan yourself.

```
int64_t ret = join("1.1=2.2&2.3=3.2");
```

query plan A



query plan B



Either way is fine.

# Tip

- We will test your project by
  1. Calling *open\_table* API multiple times to open all test tables.
  2. Calling join API multiple times with different queries and measure the time spent.
- You can preprocess in the *open\_table* API.
  - Scanning opened table, gathering some statistics, ... will be fine.
  - Timeout for *open\_table* call is **10 sec**.
    - It will be enough time to scan any database file on disk.
- You can use multiple threads for this project.
  - Our test machine spec will be announced later.